

# COSC4337\_\_109-Classification-Exercise-Solution

September 17, 2021

## 1 Classification Exercise - Solutions

We'll be working with some California Census Data, we'll be trying to use various features of an individual to predict what class of income they belong in ( $>50k$  or  $\leq 50k$ ).

Here is some information about the data:

Column Name

Type

Description

age

Continuous

The age of the individual

workclass

Categorical

The type of employer the individual has (government, military, private, etc.).

fnlwgt

Continuous

The number of people the census takers believe that observation represents (sample weight). This variable will not be used.

education

Categorical

The highest level of education achieved for that individual.

education\_num

Continuous

The highest level of education in numerical form.

marital\_status

Categorical

Marital status of the individual.

occupation

Categorical

The occupation of the individual.

relationship

Categorical

Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race

Categorical

White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

gender

Categorical

Female, Male.

capital\_gain

Continuous

Capital gains recorded.

capital\_loss

Continuous

Capital Losses recorded.

hours\_per\_week

Continuous

Hours worked per week.

native\_country

Categorical

Country of origin of the individual.

income

Categorical

“>50K” or “<=50K”, meaning whether the person makes more than \$50,000 annually.

## **1.1 Follow the Directions in Bold. If you get stuck, check out the solutions lecture.**

### **1.1.1 THE DATA**

**\*\* Read in the census\_data.csv data with pandas\*\***

```
[1]: import pandas as pd
```

```
[2]: census = pd.read_csv("census_data.csv")
```

```
[3]: census.head()
```

```
[3]:
```

	age	workclass	education	education_num	marital_status	\
0	39	State-gov	Bachelors	13	Never-married	
1	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	
2	38	Private	HS-grad	9	Divorced	
3	53	Private	11th	7	Married-civ-spouse	
4	28	Private	Bachelors	13	Married-civ-spouse	

  

	occupation	relationship	race	gender	capital_gain	\
0	Adm-clerical	Not-in-family	White	Male	2174	
1	Exec-managerial	Husband	White	Male	0	
2	Handlers-cleaners	Not-in-family	White	Male	0	
3	Handlers-cleaners	Husband	Black	Male	0	
4	Prof-specialty	Wife	Black	Female	0	

  

	capital_loss	hours_per_week	native_country	income_bracket
0	0	40	United-States	<=50K
1	0	13	United-States	<=50K
2	0	40	United-States	<=50K
3	0	40	United-States	<=50K
4	0	40	Cuba	<=50K

\*\* TensorFlow won't be able to understand strings as labels, you'll need to use pandas .apply() method to apply a custom function that converts them to 0s and 1s. This might be hard if you aren't very familiar with pandas, so feel free to take a peek at the solutions for this part.\*\*

\*\* Convert the Label column to 0s and 1s instead of strings.\*\*

```
[4]: census['income_bracket'].unique()
```

```
[4]: array([' <=50K', ' >50K'], dtype=object)
```

```
[5]: def label_fix(label):  
    if label==' <=50K':  
        return 0  
    else:  
        return 1
```

```
[6]: census['income_bracket'] = census['income_bracket'].apply(label_fix)
```

```
[7]: # Cool Alternative  
# lambda label:int(label==' <=50k')  
  
# census['income_bracket'].apply(lambda label: int(label==' <=50K'))
```

### 1.1.2 Perform a Train Test Split on the Data

```
[8]: from sklearn.model_selection import train_test_split

[9]: x_data = census.drop('income_bracket',axis=1)
     y_labels = census['income_bracket']
     X_train, X_test, y_train, y_test = train_test_split(x_data,y_labels,test_size=0.
     ↪3,random_state=101)
```

### 1.1.3 Create the Feature Columns for tf.estimator

**\*\* Take note of categorical vs continuous values! \*\***

```
[10]: census.columns

[10]: Index(['age', 'workclass', 'education', 'education_num', 'marital_status',
           'occupation', 'relationship', 'race', 'gender', 'capital_gain',
           'capital_loss', 'hours_per_week', 'native_country', 'income_bracket'],
          dtype='object')
```

**\*\* Import Tensorflow \*\***

```
[11]: import tensorflow as tf
```

**\*\* Create the tf.feature\_columns for the categorical values. Use vocabulary lists or just use hash buckets. \*\***

```
[12]: gender = tf.feature_column.categorical_column_with_vocabulary_list("gender",
     ↪["Female", "Male"])
     occupation = tf.feature_column.
     ↪categorical_column_with_hash_bucket("occupation", hash_bucket_size=1000)
     marital_status = tf.feature_column.
     ↪categorical_column_with_hash_bucket("marital_status", hash_bucket_size=1000)
     relationship = tf.feature_column.
     ↪categorical_column_with_hash_bucket("relationship", hash_bucket_size=1000)
     education = tf.feature_column.categorical_column_with_hash_bucket("education",
     ↪hash_bucket_size=1000)
     workclass = tf.feature_column.categorical_column_with_hash_bucket("workclass",
     ↪hash_bucket_size=1000)
     native_country = tf.feature_column.
     ↪categorical_column_with_hash_bucket("native_country", hash_bucket_size=1000)
```

**\*\* Create the continuous feature\_columns for the continuous values using numeric\_column \*\***

```
[13]: age = tf.feature_column.numeric_column("age")
     education_num = tf.feature_column.numeric_column("education_num")
     capital_gain = tf.feature_column.numeric_column("capital_gain")
     capital_loss = tf.feature_column.numeric_column("capital_loss")
     hours_per_week = tf.feature_column.numeric_column("hours_per_week")
```

**\*\* Put all these variables into a single list with the variable name feat\_cols \*\***

```
[14]: feat_cols = [
    → gender, occupation, marital_status, relationship, education, workclass, native_country,
    age, education_num, capital_gain, capital_loss, hours_per_week]
```

#### 1.1.4 Create Input Function

**\*\* Batch\_size is up to you. But do make sure to shuffle!\*\***

```
[15]: input_func=tf.estimator.inputs.
    → pandas_input_fn(x=X_train,y=y_train,batch_size=100,num_epochs=None,shuffle=True)
```

Create your model with `tf.estimator` Create a `LinearClassifier`. (If you want to use a `DNNClassifier`, keep in mind you'll need to create embedded columns out of the categorical feature that use strings, check out the previous lecture on this for more info.)

```
[16]: model = tf.estimator.LinearClassifier(feature_columns=feat_cols)
```

```
INFO:tensorflow:Using default config.
WARNING:tensorflow:Using temporary folder as model directory:
C:\Users\RizkN\AppData\Local\Temp\tmpc4n5bi1j
INFO:tensorflow:Using config: {'_model_dir':
'C:\\Users\\RizkN\\AppData\\Local\\Temp\\tmpc4n5bi1j', '_tf_random_seed': None,
'_save_summary_steps': 100, '_save_checkpoints_steps': None,
'_save_checkpoints_secs': 600, '_session_config': allow_soft_placement: true
graph_options {
  rewrite_options {
    meta_optimizer_iterations: ONE
  }
}
, '_keep_checkpoint_max': 5, '_keep_checkpoint_every_n_hours': 10000,
'_log_step_count_steps': 100, '_train_distribute': None, '_device_fn': None,
'_protocol': None, '_eval_distribute': None, '_experimental_distribute': None,
'_experimental_max_worker_delay_secs': None, '_session_creation_timeout_secs':
7200, '_service': None, '_cluster_spec':
<tensorflow.python.training.server_lib.ClusterSpec object at
0x00000015A54FF6848>, '_task_type': 'worker', '_task_id': 0,
'_global_id_in_cluster': 0, '_master': '', '_evaluation_master': '',
'_is_chief': True, '_num_ps_replicas': 0, '_num_worker_replicas': 1}
```

**\*\* Train your model on the data, for at least 5000 steps. \*\***

```
[17]: model.train(input_fn=input_func,steps=5000)
```

```
WARNING:tensorflow:From C:\Users\RizkN\.conda\envs\tf1\lib\site-
packages\tensorflow_core\python\training\training_util.py:236:
Variable.initialized_value (from tensorflow.python.ops.variables) is deprecated
```

and will be removed in a future version.

Instructions for updating:

Use `Variable.read_value`. Variables in 2.X are initialized automatically both in eager and graph (inside `tf.defun`) contexts.

WARNING:tensorflow:From C:\Users\RizkN\.conda\envs\tf1\lib\site-packages\tensorflow\_estimator\python\estimator\inputs\queues\feeding\_queue\_runner.py:62: `QueueRunner.__init__` (from tensorflow.python.training.queue\_runner\_impl) is deprecated and will be removed in a future version.

Instructions for updating:

To construct input pipelines, use the ``tf.data`` module.

WARNING:tensorflow:From C:\Users\RizkN\.conda\envs\tf1\lib\site-packages\tensorflow\_estimator\python\estimator\inputs\queues\feeding\_functions.py:500: `add_queue_runner` (from tensorflow.python.training.queue\_runner\_impl) is deprecated and will be removed in a future version.

Instructions for updating:

To construct input pipelines, use the ``tf.data`` module.

INFO:tensorflow:Calling model\_fn.

WARNING:tensorflow:From C:\Users\RizkN\.conda\envs\tf1\lib\site-packages\tensorflow\_core\python\feature\_column\feature\_column\_v2.py:305: `Layer.add_variable` (from tensorflow.python.keras.engine.base\_layer) is deprecated and will be removed in a future version.

Instructions for updating:

Please use ``layer.add_weight`` method instead.

WARNING:tensorflow:From C:\Users\RizkN\.conda\envs\tf1\lib\site-packages\tensorflow\_core\python\ops\resource\_variable\_ops.py:1630: calling `BaseResourceVariable.__init__` (from tensorflow.python.ops.resource\_variable\_ops) with constraint is deprecated and will be removed in a future version.

Instructions for updating:

If using Keras pass `*_constraint` arguments to layers.

WARNING:tensorflow:From C:\Users\RizkN\.conda\envs\tf1\lib\site-packages\tensorflow\_core\python\ops\embedding\_ops.py:802: where (from tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use `tf.where` in 2.0, which has the same broadcast rule as `np.where`

WARNING:tensorflow:From C:\Users\RizkN\.conda\envs\tf1\lib\site-packages\tensorflow\_estimator\python\estimator\canned\linear.py:308: `to_float` (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use ``tf.cast`` instead.

INFO:tensorflow:Done calling model\_fn.

INFO:tensorflow>Create CheckpointSaverHook.

INFO:tensorflow:Graph was finalized.

INFO:tensorflow:Running local\_init\_op.

INFO:tensorflow:Done running local\_init\_op.

WARNING:tensorflow:From C:\Users\RizkN\.conda\envs\tf1\lib\site-packages\tensorflow\_core\python\training\monitored\_session.py:882:

start\_queue\_runners (from tensorflow.python.training.queue\_runner\_impl) is deprecated and will be removed in a future version.

Instructions for updating:

To construct input pipelines, use the `tf.data` module.

INFO:tensorflow:Saving checkpoints for 0 into

C:\Users\RizkN\AppData\Local\Temp\tmpc4n5bi1j\model.ckpt.

INFO:tensorflow:loss = 69.31474, step = 1

INFO:tensorflow:global\_step/sec: 198.15

INFO:tensorflow:loss = 2771.567, step = 101 (0.514 sec)

INFO:tensorflow:global\_step/sec: 232.797

INFO:tensorflow:loss = 23.679995, step = 201 (0.431 sec)

INFO:tensorflow:global\_step/sec: 211.435

INFO:tensorflow:loss = 172.98154, step = 301 (0.469 sec)

INFO:tensorflow:global\_step/sec: 211.865

INFO:tensorflow:loss = 754.4713, step = 401 (0.474 sec)

INFO:tensorflow:global\_step/sec: 218.318

INFO:tensorflow:loss = 40.32142, step = 501 (0.457 sec)

INFO:tensorflow:global\_step/sec: 201.596

INFO:tensorflow:loss = 231.27289, step = 601 (0.498 sec)

INFO:tensorflow:global\_step/sec: 179.855

INFO:tensorflow:loss = 262.46594, step = 701 (0.555 sec)

INFO:tensorflow:global\_step/sec: 169.935

INFO:tensorflow:loss = 264.06906, step = 801 (0.592 sec)

INFO:tensorflow:global\_step/sec: 193.182

INFO:tensorflow:loss = 190.95473, step = 901 (0.514 sec)

INFO:tensorflow:global\_step/sec: 221.709

INFO:tensorflow:loss = 40.812557, step = 1001 (0.449 sec)

INFO:tensorflow:global\_step/sec: 197.527

INFO:tensorflow:loss = 36.24775, step = 1101 (0.506 sec)

INFO:tensorflow:global\_step/sec: 278.916

INFO:tensorflow:loss = 148.35156, step = 1201 (0.361 sec)

INFO:tensorflow:global\_step/sec: 259.056

INFO:tensorflow:loss = 311.86063, step = 1301 (0.384 sec)

INFO:tensorflow:global\_step/sec: 247.239

INFO:tensorflow:loss = 210.51053, step = 1401 (0.407 sec)

INFO:tensorflow:global\_step/sec: 265.578

INFO:tensorflow:loss = 33.848213, step = 1501 (0.376 sec)

INFO:tensorflow:global\_step/sec: 257.067

INFO:tensorflow:loss = 38.28836, step = 1601 (0.386 sec)

INFO:tensorflow:global\_step/sec: 268.598

INFO:tensorflow:loss = 40.716667, step = 1701 (0.372 sec)

INFO:tensorflow:global\_step/sec: 281.689

INFO:tensorflow:loss = 35.435165, step = 1801 (0.357 sec)

INFO:tensorflow:global\_step/sec: 257.729

INFO:tensorflow:loss = 67.39629, step = 1901 (0.388 sec)

INFO:tensorflow:global\_step/sec: 210.524

INFO:tensorflow:loss = 66.776535, step = 2001 (0.477 sec)

INFO:tensorflow:global\_step/sec: 232.043

INFO:tensorflow:loss = 34.25927, step = 2101 (0.429 sec)  
INFO:tensorflow:global\_step/sec: 280.299  
INFO:tensorflow:loss = 37.18171, step = 2201 (0.358 sec)  
INFO:tensorflow:global\_step/sec: 268.089  
INFO:tensorflow:loss = 58.846634, step = 2301 (0.370 sec)  
INFO:tensorflow:global\_step/sec: 223.213  
INFO:tensorflow:loss = 39.074253, step = 2401 (0.451 sec)  
INFO:tensorflow:global\_step/sec: 223.689  
INFO:tensorflow:loss = 65.20474, step = 2501 (0.447 sec)  
INFO:tensorflow:global\_step/sec: 205.317  
INFO:tensorflow:loss = 403.8542, step = 2601 (0.489 sec)  
INFO:tensorflow:global\_step/sec: 240.57  
INFO:tensorflow:loss = 208.06223, step = 2701 (0.414 sec)  
INFO:tensorflow:global\_step/sec: 244.498  
INFO:tensorflow:loss = 33.415634, step = 2801 (0.408 sec)  
INFO:tensorflow:global\_step/sec: 271.003  
INFO:tensorflow:loss = 30.884544, step = 2901 (0.370 sec)  
INFO:tensorflow:global\_step/sec: 271.738  
INFO:tensorflow:loss = 41.776295, step = 3001 (0.364 sec)  
INFO:tensorflow:global\_step/sec: 260.984  
INFO:tensorflow:loss = 44.12464, step = 3101 (0.384 sec)  
INFO:tensorflow:global\_step/sec: 218.563  
INFO:tensorflow:loss = 37.060394, step = 3201 (0.456 sec)  
INFO:tensorflow:global\_step/sec: 241.52  
INFO:tensorflow:loss = 92.91695, step = 3301 (0.421 sec)  
INFO:tensorflow:global\_step/sec: 225.288  
INFO:tensorflow:loss = 67.31664, step = 3401 (0.437 sec)  
INFO:tensorflow:global\_step/sec: 222.692  
INFO:tensorflow:loss = 104.52739, step = 3501 (0.451 sec)  
INFO:tensorflow:global\_step/sec: 251.032  
INFO:tensorflow:loss = 82.66494, step = 3601 (0.400 sec)  
INFO:tensorflow:global\_step/sec: 226.245  
INFO:tensorflow:loss = 87.70129, step = 3701 (0.443 sec)  
INFO:tensorflow:global\_step/sec: 222.195  
INFO:tensorflow:loss = 185.81793, step = 3801 (0.447 sec)  
INFO:tensorflow:global\_step/sec: 239.813  
INFO:tensorflow:loss = 69.91282, step = 3901 (0.417 sec)  
INFO:tensorflow:global\_step/sec: 255.104  
INFO:tensorflow:loss = 33.424725, step = 4001 (0.391 sec)  
INFO:tensorflow:global\_step/sec: 207.262  
INFO:tensorflow:loss = 250.0559, step = 4101 (0.486 sec)  
INFO:tensorflow:global\_step/sec: 272.411  
INFO:tensorflow:loss = 185.93808, step = 4201 (0.367 sec)  
INFO:tensorflow:global\_step/sec: 251.254  
INFO:tensorflow:loss = 27.12293, step = 4301 (0.397 sec)  
INFO:tensorflow:global\_step/sec: 268.714  
INFO:tensorflow:loss = 90.04298, step = 4401 (0.373 sec)  
INFO:tensorflow:global\_step/sec: 267.052



```
INFO:tensorflow:loss = 59.009064, step = 4501 (0.374 sec)
INFO:tensorflow:global_step/sec: 217.623
INFO:tensorflow:loss = 96.62125, step = 4601 (0.455 sec)
INFO:tensorflow:global_step/sec: 237.141
INFO:tensorflow:loss = 66.35788, step = 4701 (0.430 sec)
INFO:tensorflow:global_step/sec: 204.107
INFO:tensorflow:loss = 69.71101, step = 4801 (0.486 sec)
INFO:tensorflow:global_step/sec: 218.113
INFO:tensorflow:loss = 76.9079, step = 4901 (0.463 sec)
INFO:tensorflow:Saving checkpoints for 5000 into
C:\Users\RizkN\AppData\Local\Temp\tmpc4n5bi1j\model.ckpt.
INFO:tensorflow:Loss for final step: 39.92393.
```

```
[17]: <tensorflow_estimator.python.estimator.canned.linear.LinearClassifier at
0x15a54ff79c8>
```

### 1.1.5 Evaluation

**\*\* Create a prediction input function. Remember to only support X\_test data and keep shuffle=False. \*\***

```
[18]: pred_fn = tf.estimator.inputs.
      ↪ pandas_input_fn(x=X_test, batch_size=len(X_test), shuffle=False)
```

**\*\* Use model.predict() and pass in your input function. This will produce a generator of predictions, which you can then transform into a list, with list() \*\***

```
[19]: predictions = list(model.predict(input_fn=pred_fn))
```

```
INFO:tensorflow:Calling model_fn.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Graph was finalized.
INFO:tensorflow:Restoring parameters from
C:\Users\RizkN\AppData\Local\Temp\tmpc4n5bi1j\model.ckpt-5000
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.
```

**\*\* Each item in your list will look like this: \*\***

```
[20]: predictions[0]
```

```
[20]: {'logits': array([-0.91882783], dtype=float32),
      'logistic': array([0.2851968], dtype=float32),
      'probabilities': array([0.7148032, 0.2851968], dtype=float32),
      'class_ids': array([0], dtype=int64),
      'classes': array([b'0'], dtype=object),
      'all_class_ids': array([0, 1]),
      'all_classes': array([b'0', b'1'], dtype=object)}
```

**\*\* Create a list of only the class\_ids key values from the prediction list of dictionaries, these are the predictions you will use to compare against the real y\_test values. \*\***

```
[21]: final_preds = []  
      for pred in predictions:  
          final_preds.append(pred['class_ids'][0])
```

```
[22]: final_preds[:10]
```

```
[22]: [0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

**\*\* Import classification\_report from sklearn.metrics and then see if you can figure out how to use it to easily get a full report of your model's performance on the test data. \*\***

```
[23]: from sklearn.metrics import classification_report
```

```
[24]: print(classification_report(y_test, final_preds))
```

	precision	recall	f1-score	support
0	0.89	0.92	0.90	7436
1	0.71	0.64	0.67	2333
accuracy			0.85	9769
macro avg	0.80	0.78	0.79	9769
weighted avg	0.85	0.85	0.85	9769

## 2 Great Job!

```
[ ]:
```