

# Resampling Methods: Bootstrap Method

## Section 5.2

Cathy Poliak, Ph.D.  
cpoliak@central.uh.edu

Department of Mathematics  
University of Houston

# Resampling Methods

- **Resampling methods** involve repeatedly drawing samples from a training set and refitting a model of interest on each sample in order to obtain additional information about the fitted model.
- Could potentially be computationally expensive, because they involve fitting the same statistical method multiple times using different subsets of the training data.
- Two most commonly used resampling methods:
  - ▶ Cross-validation - can be used to estimate the test error associated with a given statistical learning method in order to evaluate its performance.
  - ▶ Bootstrap - can be used to provide a measure of accuracy of a parameter estimate or of a given statistical learning method.

# The Bootstrap Methods

- The **bootstrap** is a widely applicable and extremely powerful statistical tool that can be used to quantify the uncertainty associated with a given estimator or statistical learning method.
- The machine learning techniques that use the bootstrap is the *tree*-based models: Bagging, Random Forest, ect.
- The power of the bootstrap lies in the fact that it can be easily applied to a wide range of statistical learning methods, including some for which a measure of variability is otherwise difficult to obtain and is not automatically output by statistical software.

# Idea of The Bootstrap

- Resample from the original data - either directly or via a fitted model - to create data sets, from which the variability of the quantities of interest can be assessed without long-winded and error-prone analytical calculations.
- This approach involves repeating the original data analysis procedure with many replicate sets of data.
- The central goal is to obtain reliable standard errors, confidence intervals, and other measures of uncertainty for a wide range of problems.
- This approach can be applied in simple problems to check the adequacy of standard measures of uncertainty, to relax assumptions, and to give quick approximate solutions.
- The basic idea of bootstrap is to make inference about an estimate (such as the sample mean or sample coefficients  $\hat{\beta}_j$ ) for a population parameter  $\theta$  (such as the population mean or coefficients  $\beta_j$ ) on sample data.

# Steps for Bootstrap

1. Get a sample from a population with sample size  $n$ .
2. Draw a sample from the original sample data **with replacement** with size  $n$ , and replicate  $B$  times, each re-sampled sample is called a **Bootstrap Sample**, and there will be totally  $B$  Bootstrap Samples.
3. Evaluate the statistic of  $\theta$  for each Bootstrap Sample, and there will be a total of  $B$  estimates of  $\theta$ .
4. Construct a **sampling distribution** with these  $B$  Bootstrap statistics and use it to make further statistical inference, such as:
  - ▶ Estimating the standard error of the statistic for  $\theta$ .
  - ▶ Obtaining a confidence interval for  $\theta$ .

# The `boot` Function in R

Performing a bootstrap analysis in R entails only two steps:

1. Create a function that computes that statistic of interest.
2. Use the `boot()` function, which is part of the `boot` library, to perform the bootstrap by repeatedly sampling observations from the data set with replacement.

## Example

A thermostat used in an electrical device is to be checked for the accuracy of its design setting of 200°F. Ten thermostats were tested to determine their actual settings, resulting in the following data:

202.2   203.4   200.5   202.5   206.3   198.0   203.7   200.8   201.3   199.0

We wish to estimate the true median value of this thermostat.

```
> median.fun = function(dat, idx) median(dat[idx], na.rm = TRUE)
> boot.out.median = boot(data = temp, statistic = median.fun, R = 100)
> boot.out.median
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = temp, statistic = median.fun, R = 10000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	201.75	0.002145	0.8643586

# The Ideal and Reality in Statistics World

## Ideal World

- A standard error of our sample mean can be easily estimated and can find the estimated standard error.
- We assume we know or can estimate about the estimator's population.

## Real World

- Hard to know the information about the population or its distribution.
- The standard error of an estimate is hard to evaluate in general.

When the assumptions are violated, or when no formula exists for estimating standard errors, bootstrap is the powerful choice.



# Why Does the Simulation of the Bootstrap Work?

Let  $X_1, X_2, \dots, X_n$  be a random sample from a population  $P$  with cumulative distribution function  $F$ . And let  $M = g(X_1, X_2, \dots, X_n)$  be our statistic for the parameter of interest. What we desire to is to know  $\text{Var}(M)$ . We resample  $B$  times.

By the **Law of Large Numbers**:

$$\bar{m} = \frac{1}{B} \sum_{j=1}^B M_j \xrightarrow{P} E(M), \text{ as } B \rightarrow \infty$$

Where  $E(M)$  is the true mean of the statistic  $M$ .

In addition, the sample variance of these  $B$  statistics converges to the true variance of statistic  $M$  as  $B \rightarrow \infty$ .

$$s^2 = \frac{\sum_{j=1}^B (M_j - \bar{m})^2}{B - 1} \xrightarrow{P} \text{Var}(M), \text{ as } B \rightarrow \infty$$

Where  $\text{Var}(M)$  is the true variance of the statistic  $M$ .

# Empirical Cumulative Distribution Function

- The ECDF (empirical cumulative distribution function)  $F_n$  is a step function with jumps  $i/n$  at observation values, where  $i$  is the number of tied observations at that value.
- Missing values are ignored.
- For observations  $x = (x_1, x_2, \dots, x_n)$ ,  $F_n$  is the fraction of observations less or equal to  $t$ , i.e.,

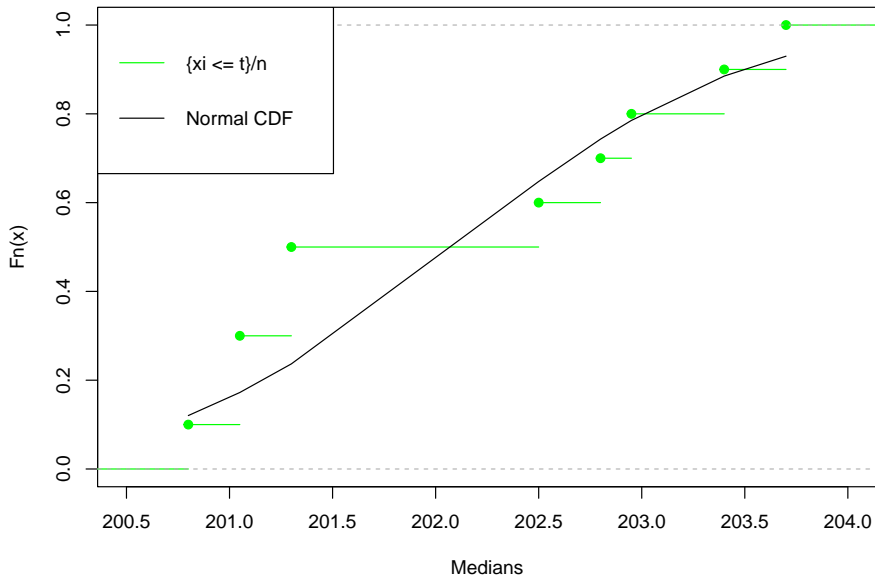
$$F_n(t) = \frac{\text{number of elements in the sample} \leq t}{n}$$

# Small Example

$$\text{median} = 206.75$$

	Samples										Medians
x1*	206.3	200.5	200.8	202.2	200.8	201.3	202.5	200.5	199	200.5	200.8
x2*	203.4	202.5	203.4	203.7	202.5	202.5	198	198	203.7	198	202.5
x3*	198	200.5	200.8	201.3	203.4	202.5	199	206.3	198	202.2	201.05
x4*	206.3	200.5	206.3	203.7	202.5	206.3	201.3	198	203.7	203.7	203.7
x5*	203.4	200.8	201.3	201.3	200.5	203.4	206.3	200.5	202.5	200.5	201.3
x6*	199	202.5	203.7	198	200.8	203.4	201.3	202.2	200.8	198	201.05
x7*	198	206.3	203.7	206.3	202.2	200.8	203.4	200.5	203.7	198	202.8
x8*	201.3	203.4	206.3	202.2	203.4	202.5	203.4	203.7	200.8	203.4	203.4
x9*	203.7	200.5	203.7	206.3	203.4	199	202.2	202.5	200.8	203.4	202.95
x10*	202.5	199	202.5	201.3	206.3	200.8	203.4	200.8	199	201.3	201.3

## Empirical Cumulative Distribution Function



# Using the `boot` Function

```
(boot2 = boot(data = temp, statistic = median.fun, R = 10))  
Bootstrap Statistics :  
      original    bias      std. error  
t1*      201.75      0.12      0.9905666  
sort(boot2$t)  
[1] 200.50 200.90 201.05 201.05 201.50 202.35 202.50 202.50 202.95 203.40
```

# Confidence Intervals

```
> #Normal  
> 201.75 + 0.032815 + c(-1,1)*qnorm(1.95/2)*sd(boot.out.median$t)  
[1] 200.0801 203.4856  
> #Percentile  
> quantile(boot.out.median$t,0.025)  
2.5%  
199.9  
> quantile(boot.out.median$t,.975)  
97.5%  
203.4  
> #Baisc  
> 2*mean(boot.out.median$t) - quantile(boot.out.median$t,0.975)  
97.5%  
200.0823  
> 2*mean(boot.out.median$t) - quantile(boot.out.median$t,0.025)  
2.5%  
203.5823
```

# Confidence Intervals

```
(ci.median = boot.ci(boot.out.median))  
BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS  
Based on 10000 bootstrap replicates
```

```
CALL :  
boot.ci(boot.out = boot.out.median)
```

```
Intervals :  
Level           Normal                Basic  
95%    (200.1, 203.4 )    (200.1, 203.6 )
```

```
Level           Percentile            BCa  
95%    (199.9, 203.4 )    (199.8, 203.4 )  
Calculations and Intervals on Original Scale
```

**For more information see:** <https://www.r-bloggers.com/2019/09/>

[understanding-bootstrap-confidence-interval-output-from-the-r-boot-package](#)

# Examples

This lecture will illustrate the use of the bootstrap in two examples

- We illustrate the bootstrap on a toy example in which we wish to determine the best investment allocation under a simple model.
- An example involving estimating the accuracy of the linear regression model.

Lab questions are in **red**. These are not multiple choice questions. Just type in your answers in blackboard.



## Example 1

Suppose that we wish to invest a fixed sum of money in two financial assets that yield returns of  $X$  and  $Y$ , respectively, where  $X$  and  $Y$  are random quantities. We will invest a fraction  $\alpha$  of our money in  $X$ , and will invest the remaining  $1 - \alpha$  in  $Y$ . Since there is variability associated with the returns on these two assets, we wish to choose  $\alpha$  to minimize the total risk, or variance, of our investment. In other words, we want to minimize  $\text{Var}(\alpha X + (1 - \alpha)Y)$ . One can show that the value that minimizes the risk is given by

$$\alpha = \frac{\sigma_Y^2 - \sigma_{XY}}{\sigma_X^2 + \sigma_Y^2 - 2\sigma_{XY}}$$

Where,  $\sigma_X^2 = \text{Var}(X)$ ,  $\sigma_Y^2 = \text{Var}(Y)$ , and  $\sigma_{XY} = \text{Cov}(X, Y)$ .

$$\text{Var}(\alpha X + (1 - \alpha)Y) = \alpha^2 \text{Var}(X) + (1 - \alpha)^2 \text{Var}(Y)$$

$$= \alpha^2 (\sigma_X^2 + \alpha(1 - \alpha) \text{Cov}(X, Y)) + (1 - \alpha)^2 (\sigma_Y^2 - \alpha(1 - \alpha) \text{Cov}(X, Y))$$

# The Estimate

In reality the population variances and covariance is unknown so we have to use estimates, using a data set that contains past measurements for  $X$  and  $Y$ . We can then estimate the value of  $\alpha$  that minimizes the variance of our investment using

$$\hat{\alpha} = \frac{\hat{\sigma}_Y^2 - \hat{\sigma}_{XY}}{\hat{\sigma}_X^2 + \hat{\sigma}_Y^2 - 2\hat{\sigma}_{XY}}$$

# Using R

1. Install the `ISLR`, we will be using the `Portfolio` data set.

```
install.packages("ISLR")  
library(ISLR)
```

2. Create the function which takes as input  $(X, Y)$  data as well as a vector indicating which observations should be used to estimate  $\alpha$ . The function then outputs the estimate for  $\alpha$  based on the selected observations. The function is as follows:

```
alpha.fn = function(data, index) {  
  X = data$X[index]  
  Y = data$Y[index]  
  return((var(Y) - cov(X, Y)) / (var(X) + var(Y) - 2*cov(X, Y)))  
}
```

Be careful about capitalization and lower case in these variables.

3. This function *returns* or outputs an estimate for  $\alpha$  based on applying the formula to the observations indexed by the argument `index`. For instance, the following command tells R to estimate  $\alpha$  using all 100 observations.

# Lab Questions

Type and run in R

```
alpha.fn(Portfolio, 1:100)
```

**Question 1:** From this command, give an estimate of  $\alpha$ . 0.5758321 =  $\alpha$

The following command uses the `sample` function to randomly select 100 observations from the range 1 to 100, with replacement. This is equivalent to constructing a new bootstrap data set and recomputing  $\alpha$  based on the new data set.

```
set.seed(10)
alpha.fn(Portfolio, sample(100, 100, replace = TRUE))
```

**Question 2:** From this command, give an estimate of  $\alpha$ .

 $\alpha$  = 0.508921

# Using the `boot` Function

We can implement a bootstrap analysis by performing this command many times, recording all of the corresponding estimate for  $\alpha$ , and computing the resulting standard deviation. However, the `boot()` function automates this approach. Below, is the function to produce  $R = 1000$  bootstrap estimates for  $\alpha$ .

```
install.packages("boot")  
library(boot)  
alpha.boot = boot(Portfolio,alpha.fn,R = 1000)  
alpha.boot  
mean(alpha.boot$t)
```

**Question 3:** What is the original estimate of  $\alpha$ ? What is  $SE(\hat{\alpha})$ ? What is the mean of the bootstrap estimates of  $\alpha$ ? *see next slide*

**Question 4:** Do the command `sd(alpha.boot$t)` compare that to the standard error from the `boot()` function. What does the standard error mean?

Answer

Bootstrap Statistics :

	original	bias	std. error
t1*	0.5758321	0.004633468	0.09184024

Question 3:

original estimate =  $\hat{\alpha} = 0.5758$

$$SE(\hat{\alpha}) = 0.0918$$

$$\text{mean}(\hat{\alpha}) = 0.5758 + 0.0046 = 0.5804$$

Question 4:

$$SE(\hat{\alpha}) = SD(t^*)$$

This is measuring the spread from the center

## Example 2

The bootstrap approach can be used to assess the variability of the coefficient estimates and predictions from a statistical learning method. We use the bootstrap approach in order to assess the variability of the estimates for  $\beta_0$  and  $\beta_1$ , the intercept and slope terms for the linear regression model that uses `horsepower` to predict `mpg` in the `Auto` data set.

## Input in R

Create a function called `boot.fn()` which takes into the `Auto` data set as well as a set of indices for the observations, and returns the intercept and slope estimates for the linear regression model.

```
boot.fn = function(data, index)
  return(coef(lm(mpg~horsepower, data = data, subset = index))
boot.fn(Auto, 1:392)
```

*Note:* We do not need `{` and `}` at the beginning and end of the defined function because it is only one line.

**Question 5:** Give the estimates of the intercept and slope.

```
> boot.fn(Auto, 1:392)
(Intercept) horsepower
39.9358610 -0.1578447
```

$$\hat{\beta}_0 = 39.9358, \quad \hat{\beta}_1 = -0.1578$$



# Lab Questions

See next slide for answers

- The `boot.fn()` function can also be used in order to create bootstrap estimates for the intercept on slope terms by randomly sampling from among the observations.

**Question 6:** Run the following command twice give two of the bootstrap estimates for the intercept and the slope.

```
boot.fn(Auto, sample(392, 392, replace = TRUE))
```

- We can use the `boot()` function to compute the standard errors of 1000 bootstrap estimates for the intercept and slope terms.

```
boot.out = boot(Auto, boot.fn, 1000)
```

```
boot.out
```

**Question 7:** What is the standard errors of the bootstrap estimates for the intercept and the slope?

- The command below gives the estimates of the original data using the `lm()` function.

```
summary(lm(mpg~horsepower, data = Auto))$coef
```

**Question 8:** Compare the standard errors given in the summary from the command above and the bootstrap function. Describe why there is a difference in the standard errors.

$$SE(\hat{\beta}_i) = \sqrt{\frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{(n-2) \sum_{i=1}^n (x_i - \bar{x})^2}}$$

The SE from the bootstrap is the standard deviation of the estimates  
The SE from 'summary' is based on the original data

## Question 6:

```
> boot.fn(Auto,sample(392,392,replace = TRUE))  
(Intercept) horsepower  
40.1186906 -0.1577063  
> boot.fn(Auto,sample(392,392,replace = TRUE))  
(Intercept) horsepower  
40.1834549 -0.1585993  
  
> boot.out = boot(Auto,boot.fn,1000)  
> boot.out
```

## ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = Auto, statistic = boot.fn, R = 1000)
```

Bootstrap Statistics :

	original	bias	std. error	
t1*	39.9358610	0.0522971635	0.84001614	$= \hat{\beta}_0$ (Intercept)
t2*	-0.1578447	-0.0006027039	0.00733243	$= \hat{\beta}_1$ (slope)

## Question 7:

$$SE(\hat{\beta}_0) = 0.84$$
$$SE(\hat{\beta}_1) = 0.007$$

NOTE: You will get slightly different answers. So your answer SHOULD NOT be exactly the same as mine.